

# Cloud Computing for **Beginners**

## Database Technologies

By Idan Gabrieli



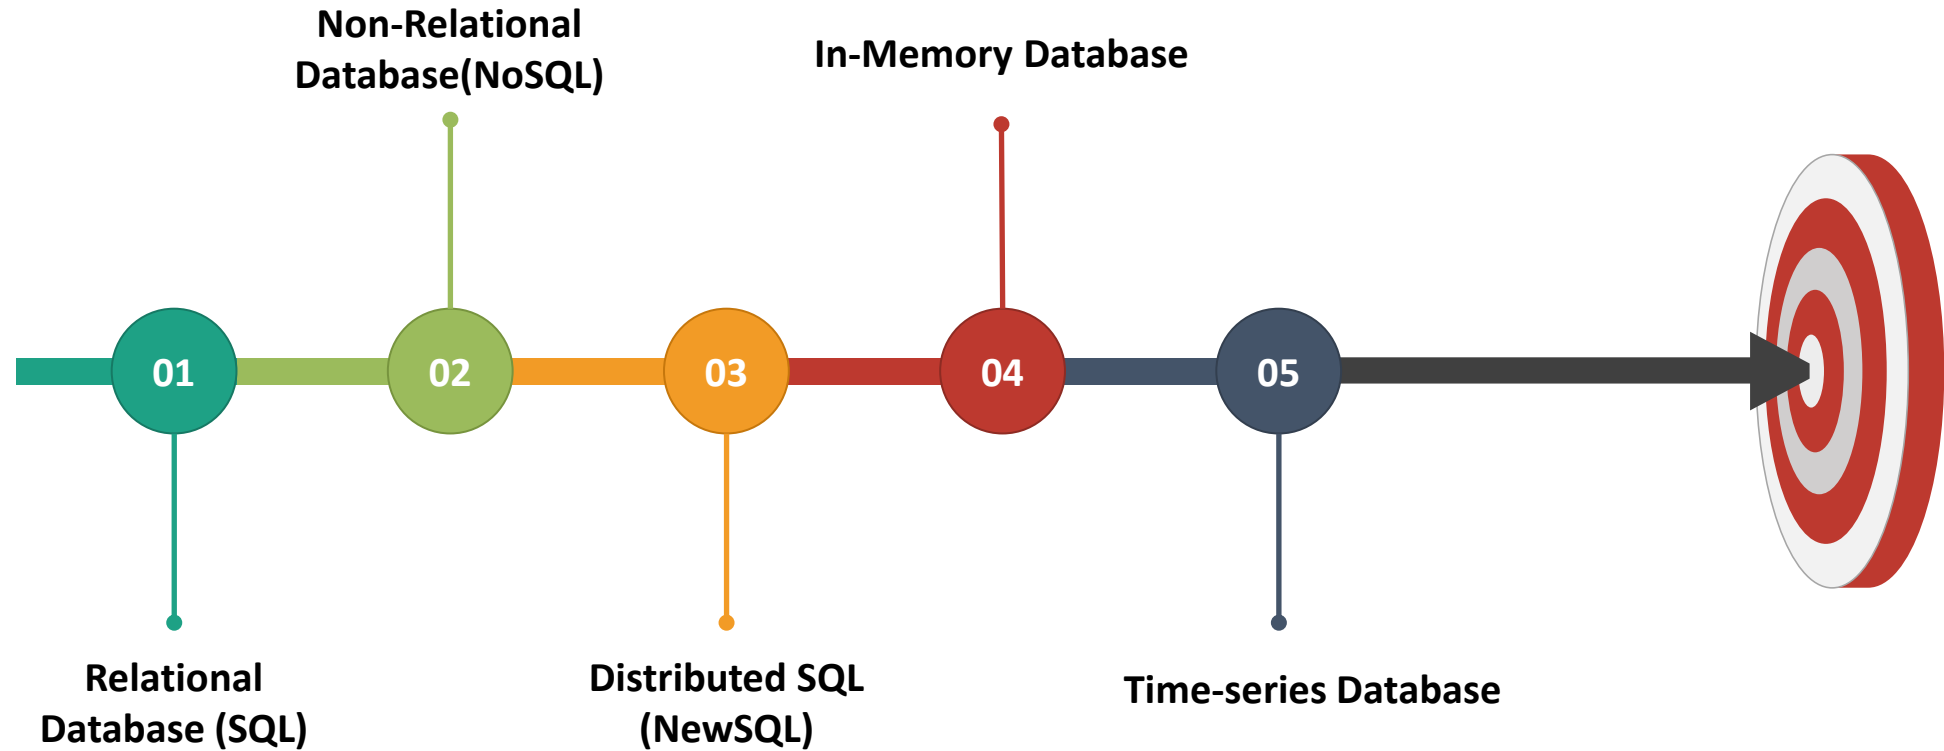


# Database Technologies

What, Where, and Why?




# Database Technologies



# Traditional Relational Database (SQL)

- A **relational** database stores data **in tables**
  - Each table is looking like an excel sheet with rows and columns
  - Each database will have a **predefined schema**
  - One or more columns in each table are used as the **primary key**
  - Rows in a table can be linked to rows in other tables using **foreign keys**



# Traditional Relational Database (SQL)

A Database Online Shop

Customers



Customer ID, Name, Phone, Address...

primary key

Products



primary key

Product Num, Description, Price, Inventory.....

Orders



primary key

Order Num, Price, Date, Customer ID, Product Num

Foreign key

Foreign key

# Traditional Relational Database (SQL)

- **SQL** - Structured Query Language
- Relational databases are a very popular option, specifically while handling operational data (**transactions**)
- The **predefined database** schema model is helping to avoid errors and keep things organized
- Popular SQL databases:



# Traditional Relational Database (SQL)

Who is using relational databases?

- **Enterprise Companies**

- Many types of applications installed on-premise/private cloud
- The dominate database technology is **relational database**

- **Internet Web Players**

- Serving **millions of customers** worldwide
- Having more challenges with scalability issues
- Websites are expected to be **on-line and available all the time**
- Traditional relational databases are NOT a great fit those such requirements



# Traditional Relational Database (SQL)


What's the Problem?

- **Scalability!**
  - One of the biggest challenges with traditional relational databases is **scalability**
  - Only **vertical scaling** (scale up/down)
    - Bigger and bigger machines
    - Unacceptable downtime when scaling up
- For web applications **availability** and **performance** are critical requirements
  - Looking for alternatives
  - Moving to non-relational databases...



# Non-relational Database (NoSQL)

- **Amazon, Facebook, Google....**
  - **BIG** internet players with **BIG** data challenges
  - Scaling up a database was an **inefficient and expensive**
- **Non-relational Databases**
  - Supporting web-based applications
  - Let's focus on two main things: **scale-out** and **Availability**



# Non-relational Database (NoSQL)

- Commercial and open-source non-relational databases



# Non-relational Database (NoSQL)

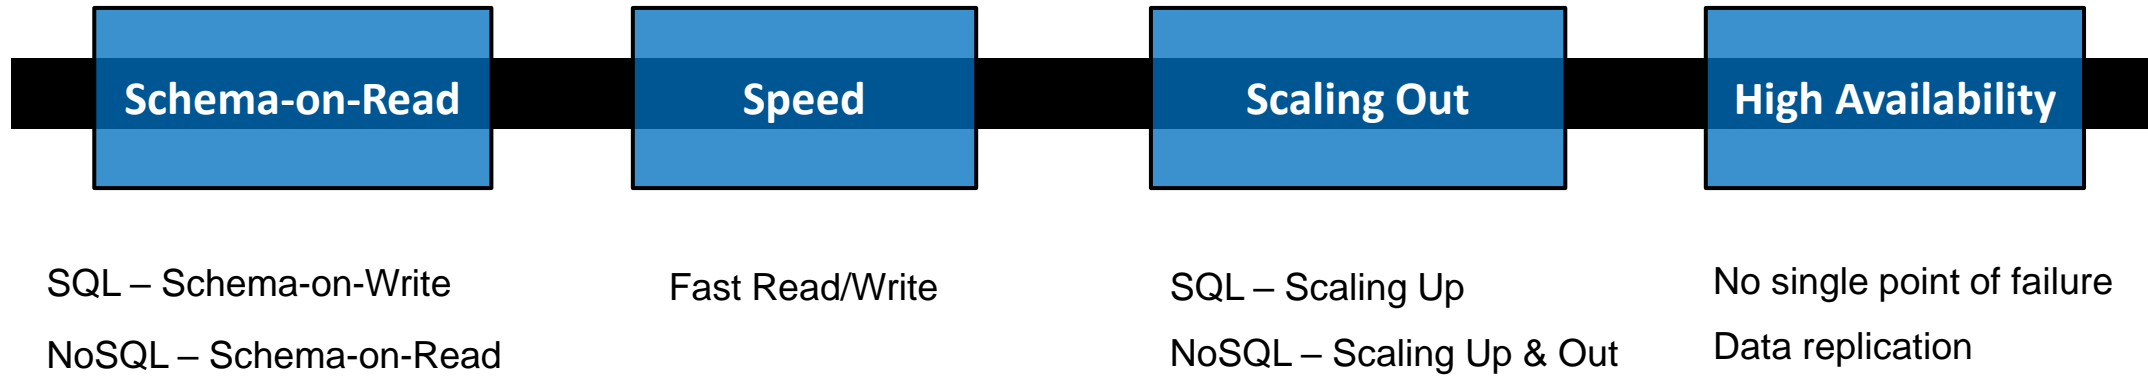
What is the meaning of a **non-relational** database?

- **Simple Data Models**
  - No predefined strong schema model
  - Using more simple data models when storing data in the database
- **NoSQL Database**
  - "No support for SQL" and then evolved to "**Not only SQL**"
- **Types of NoSQL Databases**
  - Key-value
  - Document
  - Wide column
  - Graph

# Non-relational Database (NoSQL)

What is the meaning of a non-relational database?

- Advantages





# Non-relational Database (NoSQL)

What is the meaning of a non-relational database?

- **Disadvantages**

- **Not Transactional !!!**


- Not comply with ACID model

- **Eventual Consistency Model**

- “Weaker” consistency model
    - All DB cluster will be synced with the SAME replica **AFTER** some **convergence time**
    - Reading data may not be consistent

- **NoSQL Use Cases**

- Mainly used for data analytics



# SQL and NoSQL Database


- **Traditional SQL Database**
  - **Limited scalability** and a **low resilience** to failures
  - Can handle **transactions!**
- **NoSQL Databases**
  - Designed to **scale out** using a distributed architecture
  - **Availability** is more important than **consistency**
  - Limited support for handling transactions
  - Not optimized for SQL queries
- **You win something, you lose something...**

*We believe it is better to have application programmers deal with performance problems due to overuse of transactions as bottlenecks arise, rather than always coding around the lack of transactions. (Google)*

# Distributed SQL Database (NewSQL)


- NewSQL databases emerged to make SQL scalable

*NewSQL is a class of **relational** database management systems that **seek to provide the scalability of NoSQL** systems for **online transaction processing (OLTP)** workloads while maintaining the **ACID guarantees** of a traditional database system. (Wikipedia)*



# Distributed SQL Database (NewSQL)


- A **single logical** relational database instance that is **distributed** on a global scale over **multiple geographic regions**
- Data will be **replicated** into multiple nodes inside the cluster
- **Scale out** the cluster and add more nodes





# Distributed SQL Database (**NewSQL**)

- **NewSQL Core Functionalities**
  - Supporting SQL (relational model)
  - Geographically distributed
  - Scale on demand (up and out!)
  - ACID Compliant (transactions....)
  - Deployment options
    - On premises
    - Self-hosted cloud
    - Managed service (DBaaS)



# In-memory Database (IMDB)

How the database store the data?

- **Disk-based Database**

- The most popular and traditional option
- SSD/HDD technologies
- Data is **persistence** (will not be erased after a server reboot)
- The **disk access** time is the main **bottleneck** point for databases

- **In-Memory Database (IMDB)**

- Store the data completely **in-memory**
- Can be SQL/NoSQL
- Eliminating the time needed to query data from a disk
- A memory in a computer is much more expensive than disk storage



# In-memory Database (IMDB)

How the database store the data?


- **In-Memory Database (IMDB)**
  - A great solution for applications that require **microsecond response times**
  - One of the top use cases of an in-memory database is for **data caching**
  - Popular in-memory solutions:



# Time-Series Database (TSDB)

It's about TIME!

- One of the **fastest-growing** database categories
  - Connected objects → data, data and data
  - E.g. IoT, Monitoring IT...
- **The IoT Wave**
  - Using sensors to measure something...
  - Monitor, Analyze, and Explore
  - A typical dimension that is very important → **TIME**
    - Timestamp





# Time-Series Database (TSDB)

It's about TIME!

- **Time-Series Data**
  - A series of values where the **X-axis is time**
  - Modeling Components
    - **Timestamp**
    - **Subject**
    - **Measurement/s**
  - Regular intervals, and also fluctuating or random events
  - Typically high data volumes
- **Required Database Technology**
  - High-frequency data write
  - Time functions for developers
  - Using traditional relational database???

# Time-Series Database (TSDB)

- **Time Series Database (TSDB)**
  - **Specialized** database technology for time series data
  - More **optimized** from a performance perspective
  - Automatically store recent data in-memory for fast access and move historical data to disk storage
  - Handle **high-volume data ingestion** (writing to the database)
  - Built-in specialized **time-related functions** while using SQL

